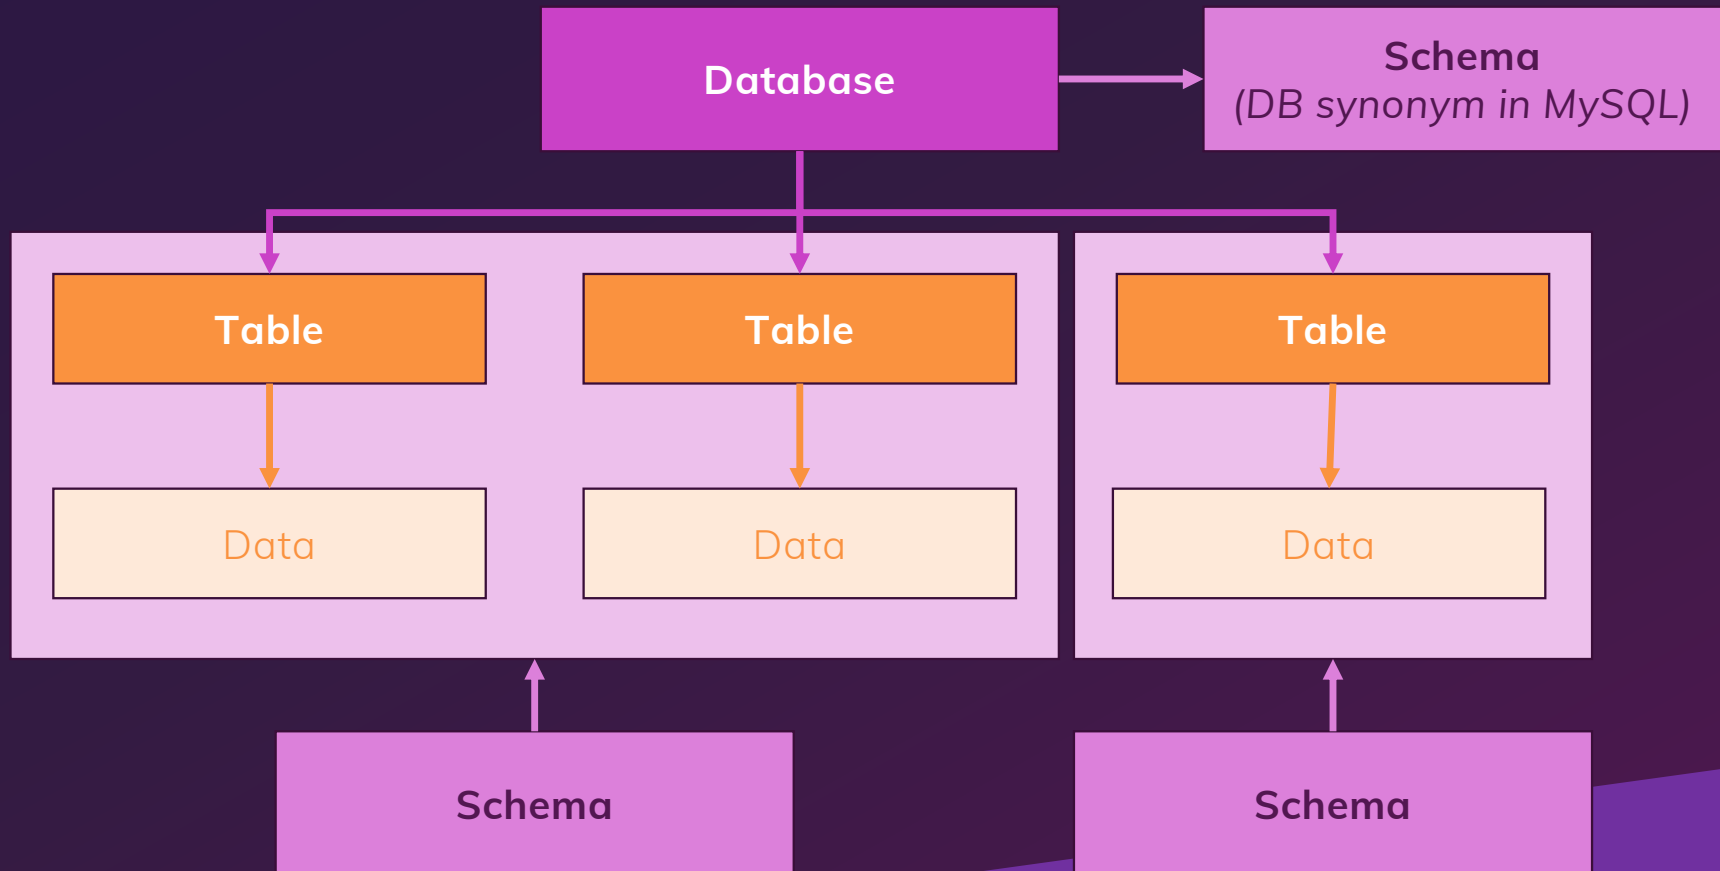
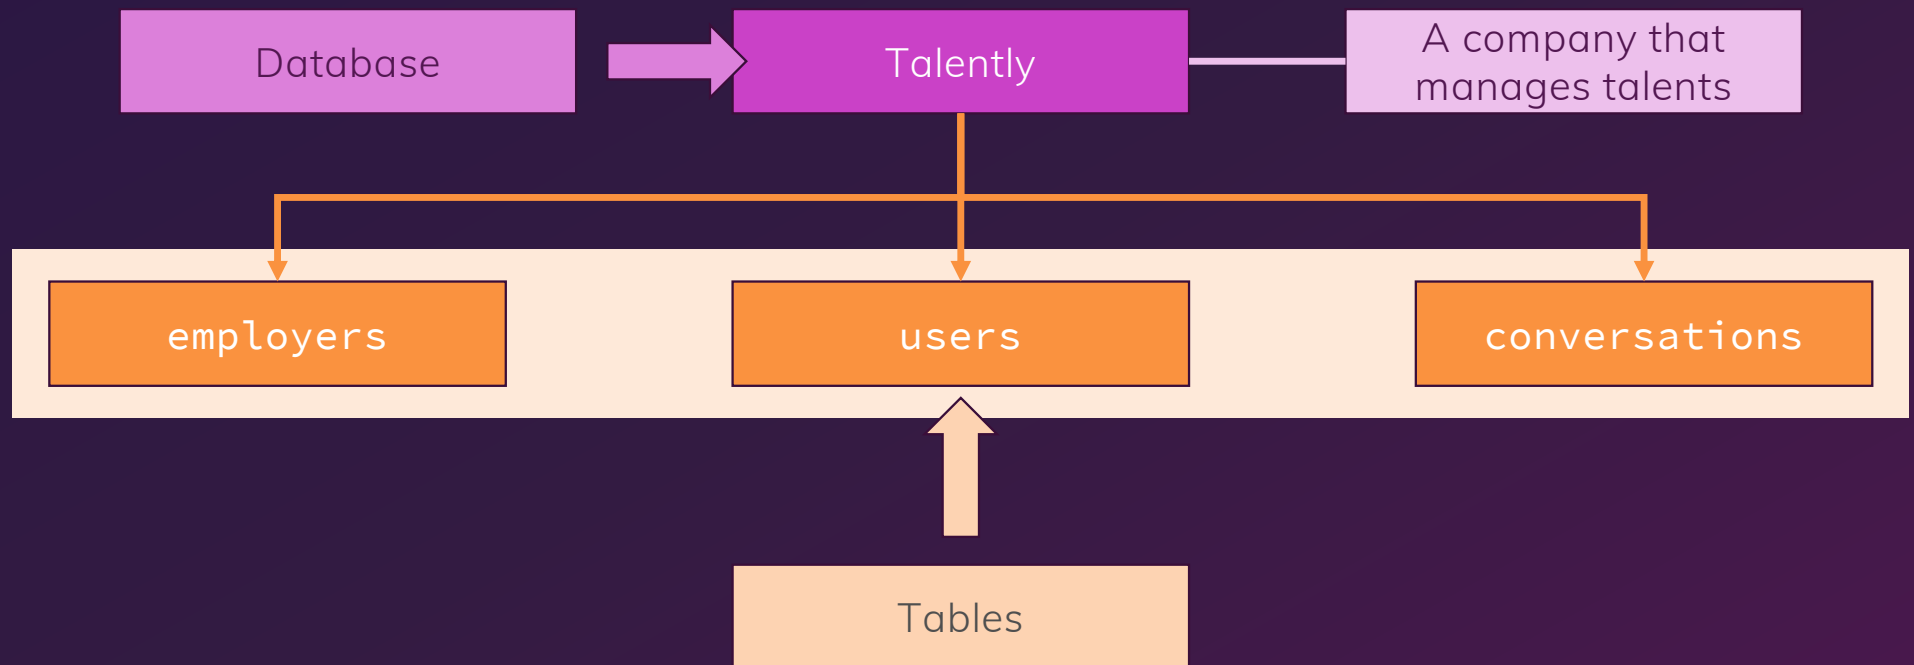


Database & Table Hierarchy



Our Example For This Course Section



Managing Databases & Tables

Data Definition Operations

Typically executed way less frequently than data manipulation operations / queries

Creating Databases

CREATE DATABASE

Configure name and base settings, e.g. regarding text handling

Creating Tables

CREATE TABLE

Configure name and fields + field value types + extra settings (e.g. constraints, keys)

Updating Databases

ALTER DATABASE

Update configuration

Updating Tables

ALTER TABLE

Update configuration

Talently: Which Data Should Be Stored?

employers

Company Name

Address (Street, City)

Yearly Revenue

Is Hiring?

users

Name (Full Name)

Yearly Salary

Status (Employed?)

Employer

conversations

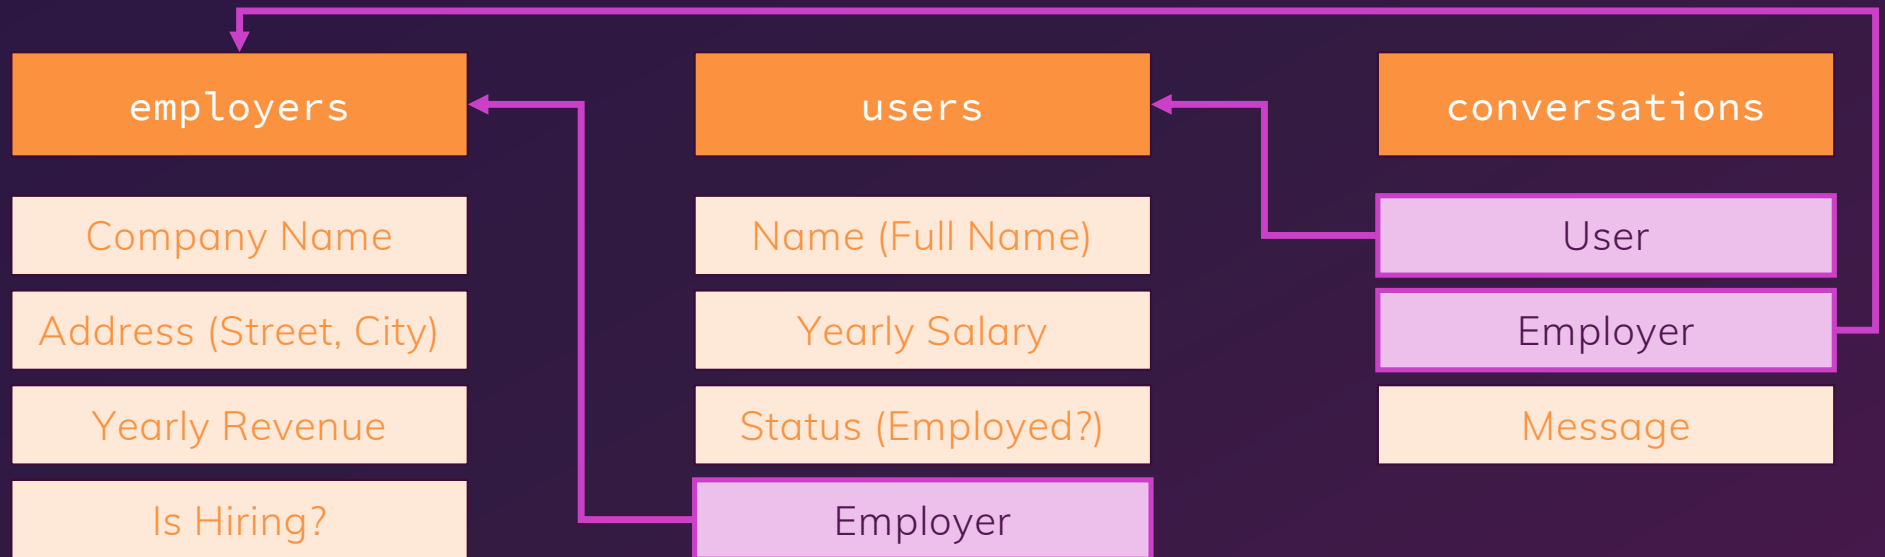
User

Employer

Message

Date Sent

Talently: Which Data Should Be Stored?



Key Data Types / Value Types

Text	Numeric	Date	Other
CHAR(X) Store text up to X characters; shorter text will be space padded	INT, SMALLINT, ... Integer numbers (between min and max boundaries) are allowed	DATE A value like 1986-10-20 (i.e. no hours or minutes)	BOOLEAN True or false (0 or 1)
VARCHAR(X) Store text up to X characters; shorter strings will not be changed	DECIMAL, NUMERIC Decimal numbers with a fixed precision (exact values)	DATETIME, TIMESTAMP A value like 1986-10-20 14:39:05 (i.e. with hours, minutes etc.)	JSON JSON-formatted text data
TEXT, LONGTEXT, ... Text of any size can be stored without specifying a max size first	FLOAT, REAL Decimal numbers with floating points (approximated values)		SERIAL An auto-incrementing integer number
ENUM Only values from a predefined set of allowed values are accepted	Not all types are part of the official standard – and not all database systems support all types		

Integer Values

5

10

-20

...

Number Values With Decimal Places

3.14

5.58

-10.999

...

CHAR vs VARCHAR vs TEXT (vs LONGTEXT ...)

Pre-defined maximum length

CHAR(X)

VARCHAR(X)

Typically used!

Text with max. length of X bytes

One byte can be one character → Depends on encoding

Shorter text is space-padded

Shorter text is not changed

CHAR(4)

VARCHAR(4)

Inserted

'hi'

Inserted

'hi'

Stored

' hi'

Stored

'hi'

No maximum length

(database system limits apply)

TEXT

LONGTEXT, ...

Typically used!

Text with no user-defined max. length (max. length depends on data type)

One byte can be one character → Depends on encoding

Max. size is 1GB in Postgres, 65,535 characters in MySQL

Not supported in Postgres, different types with different sizes in MySQL

Not part of the SQL standard but supported by many database systems

A Closer Look At Numeric Value Types

Integer ("Whole") Numbers	Exact Decimal Point Numbers	Approximate Decimal Point Numbers
INT, SMALLINT, ...	DECIMAL, NUMERIC	FLOAT, REAL, ...
3, -10, -1831, 9418125	724.12, -8.195, 51413.1	724.12, -8.195, 51413.1
Numbers without any decimal places	Numbers with decimal places and exact precision	Numbers with decimal places and approximate precision
Inserted numbers with decimal places are rounded	Inserted numbers are stored exactly (no data loss)	Stored approximately (data loss is possible)
Great for mathematical calculations	Great for data that requires exactness (e.g. monetary)	Great for numeric data where exactness is not required
Great performance	Slow performance	Great performance
Different types of integers occupy different amounts of space	Precision can be set when the table is created	Different types of numbers occupy different amounts of space

Storing The Salary

Storing The Salary Exactly

19,000.12

Stored exactly, so
that it can be used in
calculations without
data loss /
inaccuracy

Comes at a
performance penalty

Storing The Salary
Approximately

19,000.12 vs
19,000.13

Stored approximately
as exact values
might not matter

Good performance!

Storing The Salary As An
Integer ("Whole") Number

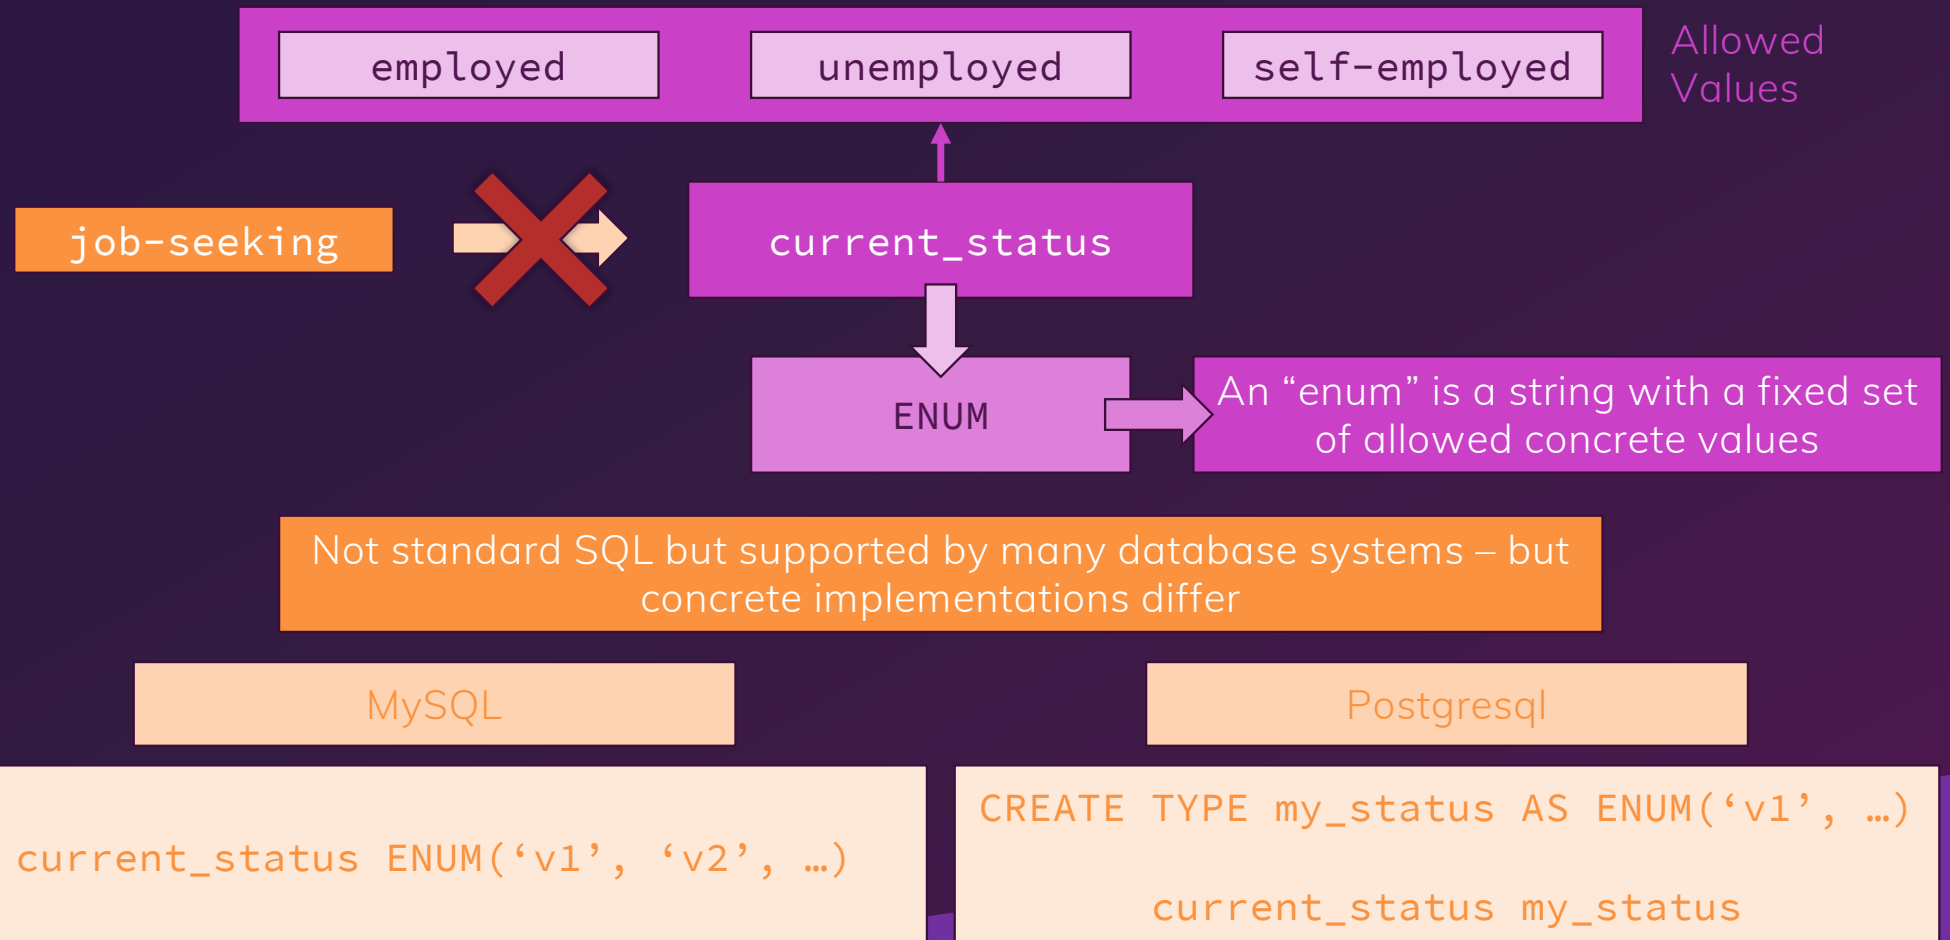
19,000

The decimal values
might not matter at
all when talking
about salaries

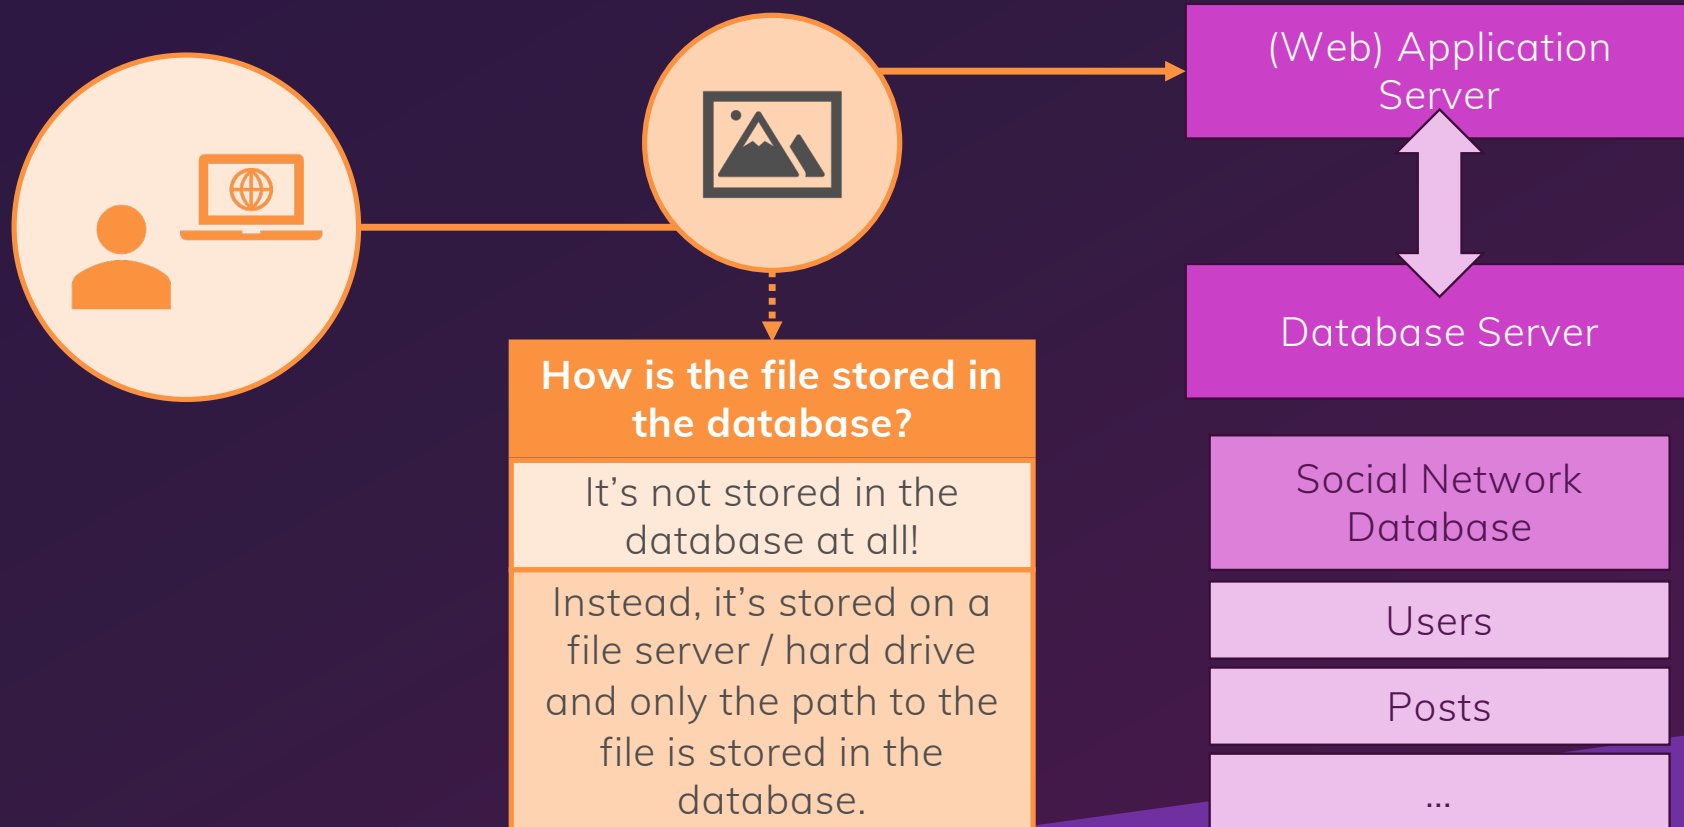
Good performance!

Floating Point Numbers / Fractional Numbers Are Challenging For Computers

Enums



What About Files?



(Not) Storing Files In Databases

```
CREATE TABLE users (  
    user_name VARCHAR(255),  
    image_path VARCHAR(200),  
);
```

A column that should hold the image file paths

```
INSERT INTO users (user_name, image_path)  
VALUES ('DBMax', 'uploads/images/db_max.jpg');
```

The path to the image (which is stored on some file server)

The Problem With “No Data”

first_name	last_name	salary	status
Max	Schwarz	15000	self-employed
Julie	Barnes	19000	employed
Michael	Smith	0	unemployed

↑
Might distort
analyses

Average: 11,333.33

The NULL Value

first_name	last_name	salary	status
Max	Schwarz	15000	self-employed
Julie	Barnes	19000	employed
Michael	Smith	[NULL]	unemployed

Shouldn't be
null!

Is ignored

Average: 17,000

Allowing Or Forbidding NULL Values

```
CREATE TABLE users (  
    full_name VARCHAR(255) NOT NULL,  
    salary INT -- NULL is allowed because it's not forbidden  
);
```

NOT NULL is a "Constraint"

This column must contain a (valid)
value – omitting it is not possible

What If Multiple Users Have The Same Name?

The Role & Importance Of Unique IDs

When storing data, each data entry should have **at least one unique value** (for identifying the record)

Unique ID

user1

user2

...

Choose & set IDs manually

hefzar32qjka

nlonkj147rkn

...

Generate unique, random strings automatically

1

2

...

Generate auto incrementing integers

Popular choice!

Setting Unique IDs & Primary Keys

```
CREATE TABLE users (  
  id INT NOT NULL UNIQUE,  
  full_name VARCHAR(255) NOT NULL,  
  salary INT  
);
```

UNIQUE is a "Constraint"

This column must not contain
duplicate values

Setting Unique IDs & Primary Keys

```
CREATE TABLE users (  
  id INT PRIMARY KEY AUTO_INCREMENT,  
  full_name VARCHAR(255) NOT NULL,  
  salary INT  
);
```

AUTO_INCREMENT is an
"Attribute"

The db system will automatically
insert an incrementing value

PRIMARY KEY is a "Constraint"

This column must be unique and
NOT NULL; Only one PRIMARY
KEY per table is allowed!



In most SQL databases / environments
(e.g. Postgresql), AUTO_INCREMENT is
not supported.

Use SERIAL PRIMARY KEY instead of
INT PRIMARY KEY AUTO_INCREMENT

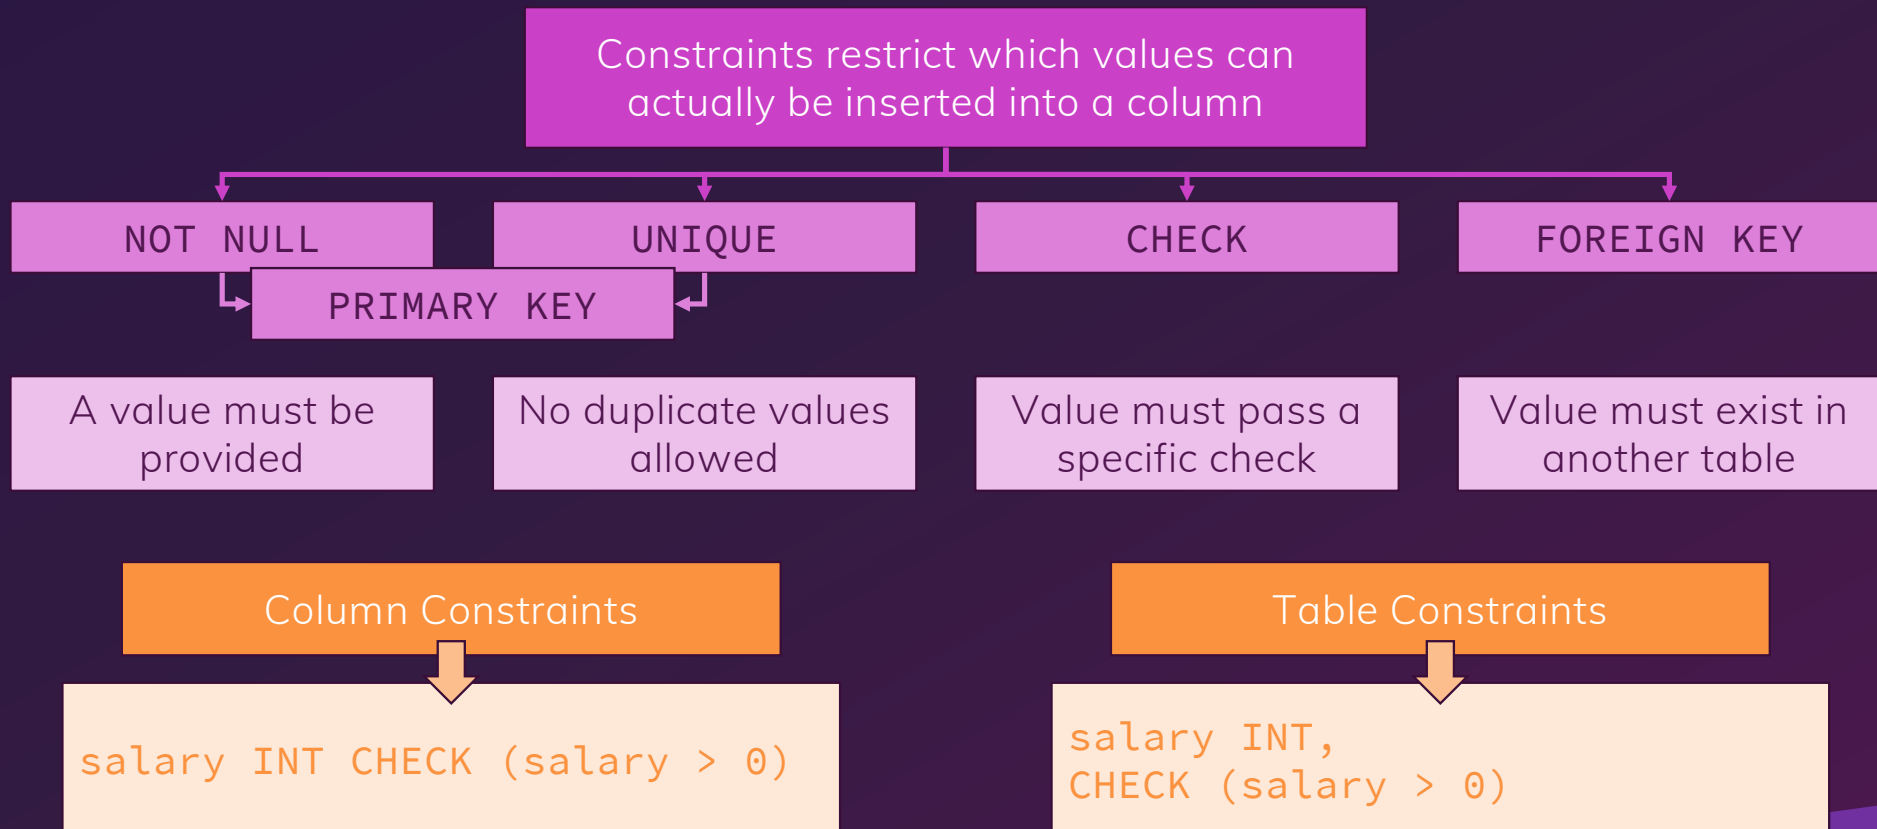
Setting Unique IDs & Primary Keys (Postgres)

```
CREATE TABLE users (  
  id SERIAL PRIMARY KEY,  
  full_name VARCHAR(255) NOT NULL,  
  salary INT  
);
```

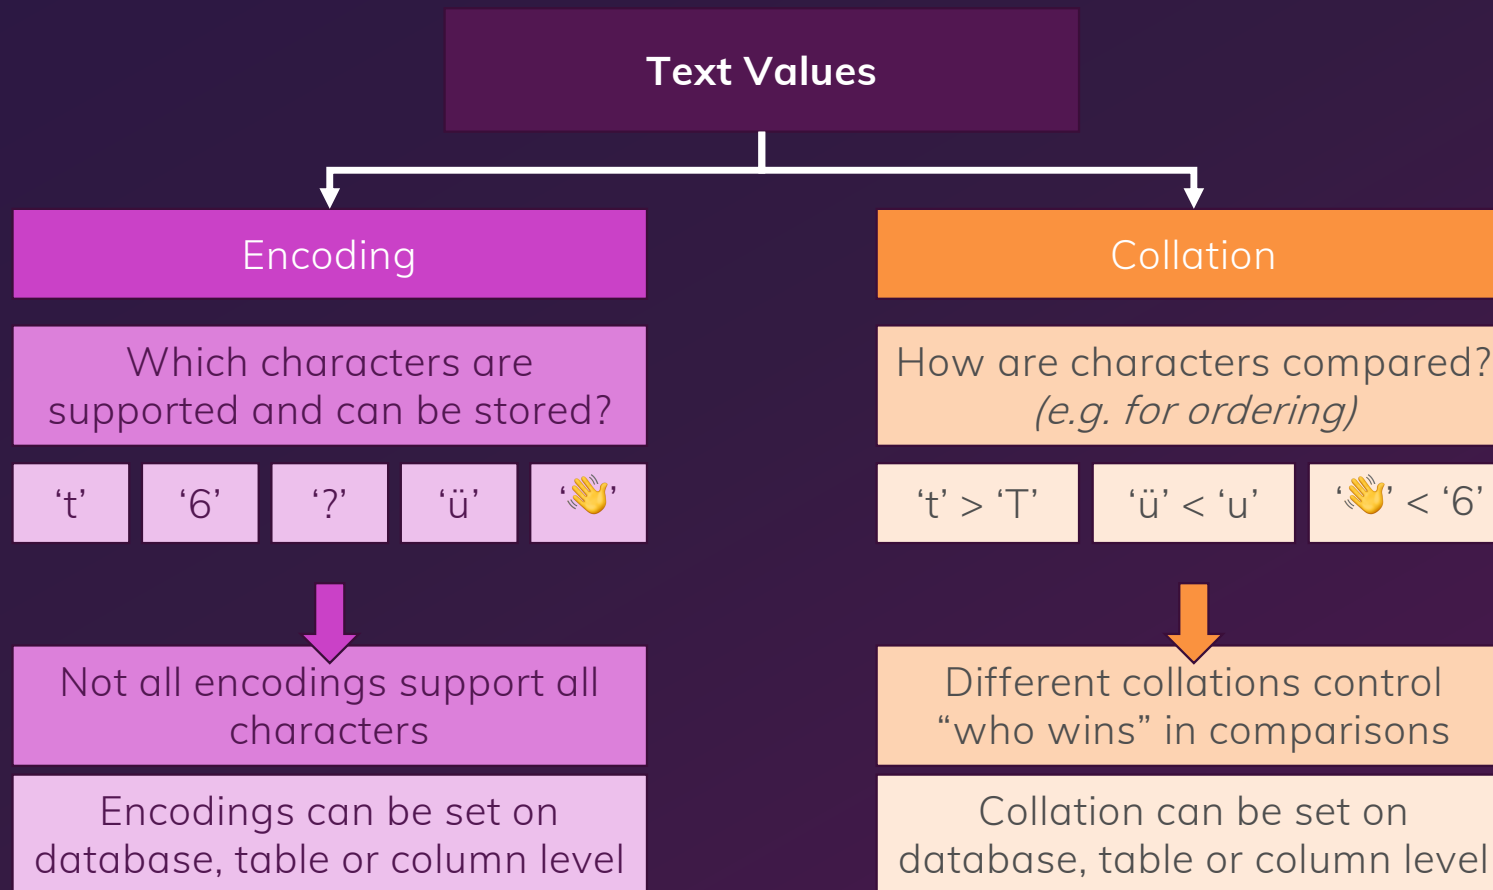
SERIAL is a Special Data Type

Creates an auto-incrementing
integer
(not supported in MySQL)

More On Constraints



Text: Encoding & Collation



More Ways Of Creating Tables

Temporary Tables

```
CREATE TEMPORARY TABLE ...
```



Tables that are only stored temporarily
(in memory of the database server)

Useful for non-permanent data
(e.g. intermediate results)

Tables Based On Other Tables / Data

```
CREATE TABLE ... AS <query>
```



Creates a table and pre-populates it with
data from a query result set

Useful if a subset of data from another
table should be stored in a separate table

Generated Columns

first_name	last_name	full_name
'Max'	'Schwarz'	'Max Schwarz'
'Julie'	'Barnes'	'Julie Barnes'

Generated Columns

first_name

‘Max’

‘Julie’

last_name

‘Schwarz’

‘Barnes’

full_name

‘Max Schwarz’

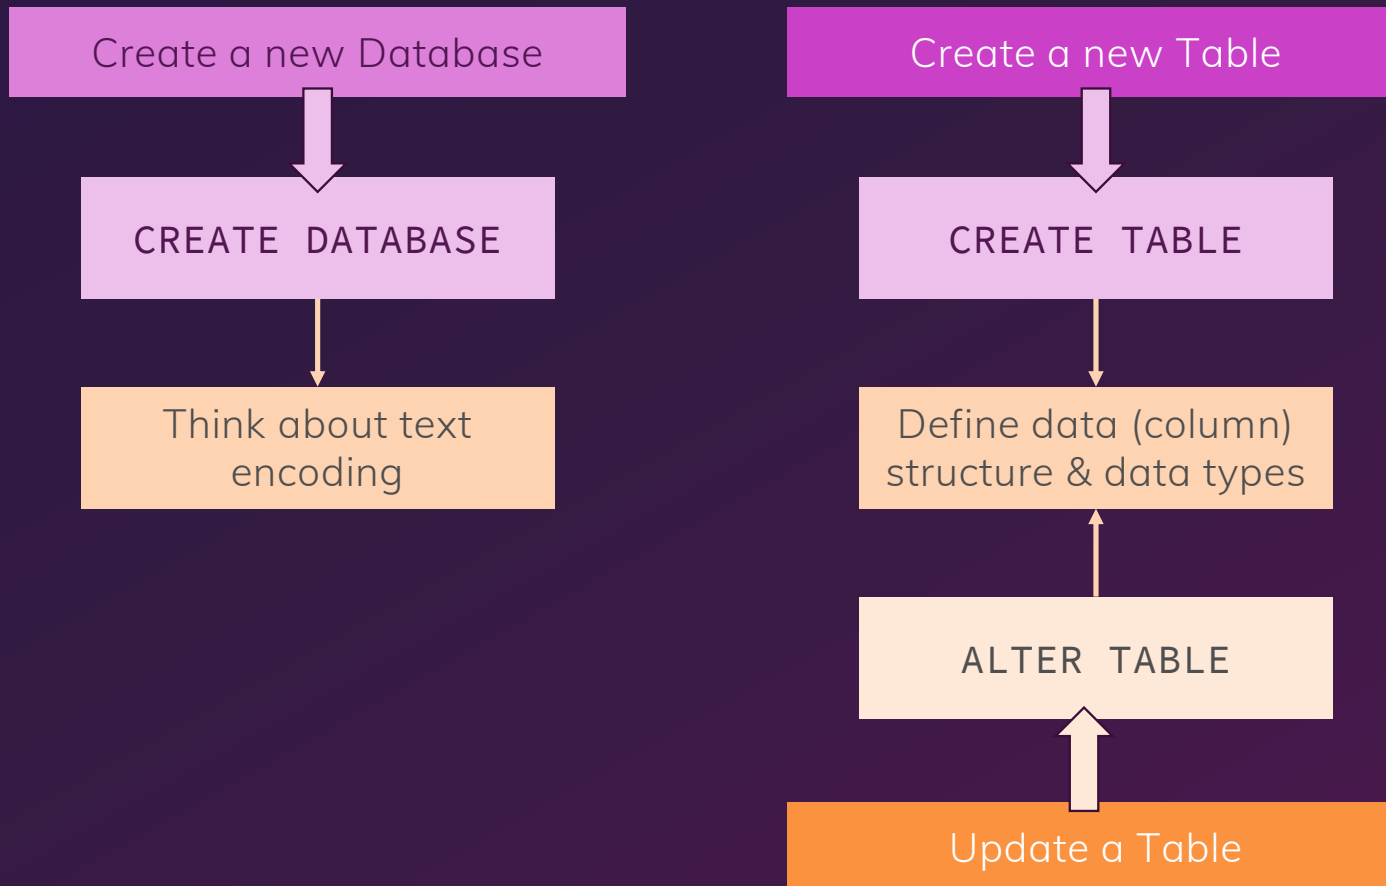
‘Julie Barnes’

Could be created as a
“Generated Column”

Value is derived automatically

Data doesn't have to be (and
can't be) inserted manually

Module Summary



Exercise Time

